

# CSC4005 Project 4 Report

---

Jiarui Chen 120090361

## Task 1: Sequential softmax implementation

In this part, a sequential softmax regression has been implemented. The output result is as follow:

```
Softmax Sequential
Training softmax regression
| Epoch | Train Loss | Train Err | Test Loss | Test Err |
| 0 | 0.35134 | 0.10182 | 0.33588 | 0.09400 |
| 1 | 0.32142 | 0.09268 | 0.31085 | 0.08730 |
| 2 | 0.30802 | 0.08795 | 0.30097 | 0.08550 |
| 3 | 0.29987 | 0.08532 | 0.29558 | 0.08370 |
| 4 | 0.29415 | 0.08323 | 0.29215 | 0.08230 |
| 5 | 0.28980 | 0.08182 | 0.28973 | 0.08090 |
| 6 | 0.28633 | 0.08085 | 0.28793 | 0.08080 |
| 7 | 0.28345 | 0.07997 | 0.28651 | 0.08040 |
| 8 | 0.28099 | 0.07923 | 0.28537 | 0.08010 |
| 9 | 0.27886 | 0.07847 | 0.28442 | 0.07970 |
Execution Time: 7002 milliseconds
```

At first, I implemented the matrix multiplication using for loops and memory locality, but at the end it takes about 17000ms. So, I use SIMD to accelerate the matrix multiplication and get the result of 7000ms.

## Task 2: OpenACC softmax

In this part, an OpenACC implementation of softmax regression has been implemented. The output result is as follow:

## Softmax OpenACC

### Training softmax regression (GPU)

Epoch	Train Loss	Train Err	Test Loss	Test Err
0	0.35134	0.10182	0.33588	0.09400
1	0.32142	0.09268	0.31085	0.08730
2	0.30802	0.08795	0.30097	0.08550
3	0.29987	0.08532	0.29558	0.08370
4	0.29415	0.08323	0.29215	0.08230
5	0.28980	0.08182	0.28973	0.08090
6	0.28633	0.08085	0.28793	0.08080
7	0.28345	0.07997	0.28651	0.08040
8	0.28099	0.07923	0.28537	0.08010
9	0.27886	0.07847	0.28442	0.07970

Execution Time: 6605 milliseconds

We use `#pragma acc data copyin` to send the data into GPU, and using `#pragma acc parallel loop present` to start parallel by OpenACC.

The implementation is based on non-SIMD sequential implementation, so it has accelerated about 10,000ms.

## Task 3: Sequential neural network implementation

In this part, a sequential nn training has been implemented. The output result is as follow:

## NN Sequential

Training two layer neural network w/ 400 hidden units

Epoch	Train Loss	Train Err	Test Loss	Test Err
0	0.13465	0.04023	0.14293	0.04240
1	0.09652	0.03020	0.11593	0.03700
2	0.07343	0.02210	0.10038	0.03190
3	0.05810	0.01693	0.09058	0.02860
4	0.04665	0.01288	0.08342	0.02650
5	0.03910	0.01025	0.07944	0.02540
6	0.03255	0.00810	0.07549	0.02440
7	0.02810	0.00687	0.07357	0.02410
8	0.02435	0.00548	0.07189	0.02340
9	0.02133	0.00463	0.07082	0.02210
10	0.01869	0.00372	0.06978	0.02200
11	0.01670	0.00310	0.06907	0.02180
12	0.01496	0.00250	0.06849	0.02150
13	0.01347	0.00193	0.06785	0.02130
14	0.01210	0.00153	0.06747	0.02060
15	0.01126	0.00125	0.06742	0.02060
16	0.01012	0.00097	0.06709	0.02040
17	0.00944	0.00078	0.06706	0.02040
18	0.00868	0.00065	0.06688	0.02040
19	0.00797	0.00045	0.06667	0.02030

Execution Time: 483065 milliseconds

The nn training implementation also use the SIMD version of matrix multiplication.

## Compile and Run

We can compile and run the programs by

```
cd ./project4
bash test.sh
```

Or run the compiled files by `sbatch` command on the cluster.

```
cd ./project4
sbatch sbatch.sh
```

## Profiling Results

We use `nsys` to get profiling results in the docker container.

The profiling files of the programs are in the `profiling` folder.

Three of the tasks which have taken most time (unit: ns) are:

IMPLEMENTS	POLL	IOCTL	READ
softmax OpenACC	3104m	128m	14.8m